

## SoME v3.5 (Decay Fix) Code

```
# Cell 1: Setup and Dependencies
print("--- Part 1: Setup and Dependencies ---")
!pip install torch datasets transformers huggingface_hub tokenizers matplotlib umap-learn
seaborn -q

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from transformers import PreTrainedTokenizerFast
from tokenizers import Tokenizer
from tokenizers.models import BPE
from tokenizers.trainers import BpeTrainer
from tokenizers.pre_tokenizers import Whitespace
from datasets import load_dataset
import copy
from tqdm import tqdm
import math
import os
import numpy as np
import matplotlib.pyplot as plt

# Verify that a GPU is available and set the device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

# Enable TF32 for A100 GPUs (Crucial for 40GB optimization)
if torch.cuda.is_available() and torch.cuda.get_device_capability(0)[0] >= 8:
    print("A100 GPU detected. Enabling TF32.")
    torch.set_float32_matmul_precision('high')

# Enable benchmark mode for cuDNN
torch.backends.cudnn.benchmark = True



---


# Cell 2: Core SoME Framework
# (Classes, Data Functions, Training Loop)

# --- 1. Model Component Classes ---

class Expert(nn.Module):
    """An expert module with configurable random weight initialization."""

```

```

def __init__(self, d_model, d_ffn, init_method='default'):
    super().__init__()
    self.w_down = nn.Linear(d_model, d_ffn)
    self.activation = nn.GELU()
    self.w_up = nn.Linear(d_ffn, d_model)

    if init_method == 'orthogonal':
        nn.init.orthogonal_(self.w_down.weight)
        nn.init.orthogonal_(self.w_up.weight)
    elif init_method == 'sparse':
        nn.init.sparse_(self.w_down.weight, sparsity=0.5)
        nn.init.sparse_(self.w_up.weight, sparsity=0.5)
    elif init_method != 'default':
        raise ValueError(f"Unknown initialization method: {init_method}")

    nn.init.zeros_(self.w_down.bias)
    nn.init.zeros_(self.w_up.bias)

def forward(self, x):
    return self.w_up(self.activation(self.w_down(x)))

class SOMELayer(nn.Module):
    def __init__(self, d_model, some_config):
        super().__init__()
        self.d_model = d_model
        self.numExperts = some_config['num_experts']
        self.d_ffn = some_config['d_ffn']
        self.top_k = some_config['top_K']

        # Heuristic update parameters
        self.alpha = some_config['alpha']
        self.beta = some_config['beta']
        self.delta = some_config['delta']

    # v4 Feature: Respawn Threshold
    self.respawn_threshold = some_config.get('respawn_threshold', 0.1)

    # Key management parameters
    self.theta_percentile = some_config['theta_percentile']
    self.warmup_steps = some_config['warmup_steps']
    self.ema_decay = some_config['ema_decay']

    self.ablation_flags = some_config.get('ablation_flags',
                                         {'use_alpha': True, 'use_beta': True, 'use_delta': True})

```

```

self.query_network = nn.Linear(d_model, d_model)

# Initialize keys on unit sphere
keys = torch.randn(self.num_experts, d_model)
self.register_buffer("key_store", F.normalize(keys, p=2, dim=-1))
self.register_buffer("usage_count", torch.zeros(self.num_experts))
self.register_buffer("steps", torch.tensor([0], dtype=torch.long))

self.experts = nn.ModuleList([Expert(d_model, self.d_ffn,
    init_method=some_config['init_method']) for _ in range(self.num_experts)])

# IN-INFERENCE PROTOCOL: Freeze all expert parameters
# We are teaching the router to use random tools.
for expert in self.experts:
    for param in expert.parameters():
        param.requires_grad = False

if self.top_k > 1:
    self.register_buffer("peer_pull_indices",
        torch.combinations(torch.arange(self.top_k), r=2))

def forward(self, x, temperature=1.0):
    batch_size, seq_len, _ = x.shape
    x_flat = x.view(-1, self.d_model)

    queries_raw = self.query_network(x_flat)
    queries = F.normalize(queries_raw, p=2, dim=-1)

    # Dot product scores.
    # v4 Fix: We rely on the magnitude of the key for "health".
    # Decayed keys will naturally have low dot products.
    scores = torch.matmul(queries, self.key_store.t())

    top_k_scores, top_k_indices = torch.topk(scores, self.top_k, dim=-1)
    gating_weights = F.softmax(top_k_scores / temperature, dim=-1)

    flat_top_k_indices = top_k_indices.view(-1)

    # Standard MoE Routing Logic
    sorted_indices, permutation_map = torch.sort(flat_top_k_indices)
    unique_expert_ids, counts = torch.unique_consecutive(sorted_indices,
    return_counts=True)

```

```

flat_inputs = x_flat.repeat_interleave(self.top_k, dim=0)
permuted_inputs = flat_inputs[permutation_map]
split_inputs = torch.split(permuted_inputs, counts.tolist(), dim=0)

output_chunks = []
for i, expert_id in enumerate(unique_expert_ids):
    output_chunks.append(self.experts[expert_id](split_inputs[i]))

concatenated_outputs = torch.cat(output_chunks, dim=0)
inverse_permutation_map = torch.argsort(permutation_map)
expert_outputs = concatenated_outputs[inverse_permutation_map]

weighted_outputs = (expert_outputs.view(-1, self.top_k, self.d_model) *
                     gating_weights.unsqueeze(-1).sum(dim=1))
final_output = weighted_outputs.view(batch_size, seq_len, self.d_model)

return x + final_output, queries, top_k_indices

@torch.no_grad()
def update_keys(self, queries, top_k_indices):
    self.steps += 1

    # Update usage counts
    unique_indices, counts = torch.unique(top_k_indices, return_counts=True)
    self.usage_count.mul_(self.ema_decay)
    self.usage_count.index_add_(0, unique_indices, (1.0 - self.ema_decay) * counts.float())

    # --- Heuristic 1: Alpha (Attraction) ---
    if self.ablation_flags.get('use_alpha', True):
        for i in range(self.top_k):
            indices = top_k_indices[:, i]
            # Calculate inertia
            inertia = 1.0 + self.usage_count[indices]
            alpha_effective = self.alpha / inertia.unsqueeze(-1)

            # Update: Move key towards query
            # Logic: k_new = k + alpha * (q - k)
            update_vec = queries - self.key_store[indices]
            self.key_store.index_add_(0, indices, alpha_effective * update_vec)

    # --- Heuristic 2: Beta (Peer Pull) ---
    if self.top_k > 1 and self.ablation_flags.get('use_beta', True):
        indices_i = top_k_indices[:, self.peer_pull_indices[:, 0]].reshape(-1)
        indices_j = top_k_indices[:, self.peer_pull_indices[:, 1]].reshape(-1)

```

```

keys_i, keys_j = self.key_store[indices_i], self.key_store[indices_j]

inertia_i = (1.0 + self.usage_count[indices_i]).unsqueeze(-1)
inertia_j = (1.0 + self.usage_count[indices_j]).unsqueeze(-1)
beta_effective = self.beta / torch.min(inertia_i, inertia_j)

update_vec_i = beta_effective * (keys_j - keys_i)
update_vec_j = beta_effective * (keys_i - keys_j)

self.key_store.index_add_(0, indices_i, update_vec_i)
self.key_store.index_add_(0, indices_j, update_vec_j)

# --- Heuristic 3: Delta (Decay) ---
# v4 Fix: We removed the global F.normalize logic.
# Now, shrinkage is permanent until the expert is used or respawned.
if self.steps > self.warmup_steps and self.ablation_flags.get('use_delta', True):
    active_usage_counts = self.usage_count[self.usage_count > 0]
    if active_usage_counts.numel() > 0:
        dynamic_theta = torch.quantile(active_usage_counts.float(), self.theta_percentile)
        low_usage_mask = self.usage_count < dynamic_theta

        # Shrink the vectors of unused experts
        self.key_store[low_usage_mask] *= (1.0 - self.delta)

# --- New: The Phoenix Mechanism (Respawn) ---
# Check for dead experts (norm < threshold) and reincarnate them
key_norms = self.key_store.norm(p=2, dim=-1)
dead_mask = key_norms < self.respawn_threshold

if dead_mask.any():
    num_dead = dead_mask.sum().item()
    # Pick random queries from the current batch to seed the new experts
    flat_queries = queries.view(-1, self.d_model)
    if flat_queries.size(0) >= num_dead:
        rand_indices = torch.randperm(flat_queries.size(0),
device=queries.device)[:num_dead]
        new_keys = flat_queries[rand_indices].clone()

        # Add slight noise and normalize to unit sphere
        new_keys = F.normalize(new_keys + torch.randn_like(new_keys) * 0.01, p=2, dim=-1)

        # Assign new keys
        self.key_store[dead_mask] = new_keys

```

```

# Reset usage stats for respawned experts (give them a chance)
self.usage_count[dead_mask] = 0.0

class SOMETwitterBlock(nn.Module):
    def __init__(self, d_model, num_heads, some_config):
        super().__init__()
        self.attention = nn.MultiheadAttention(d_model, num_heads, batch_first=True)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.some_layer = SOMELayer(d_model, some_config)

    def forward(self, x, temperature=1.0):
        seq_len = x.size(1)
        mask = torch.triu(torch.ones(seq_len, seq_len, device=x.device) * float('-inf'), diagonal=1)
        attn_output, _ = self.attention(x, x, x, attn_mask=mask)
        x = self.norm1(x + attn_output)

        # Pass temperature to the SOME layer
        some_output, queries, top_k_indices = self.some_layer(x, temperature=temperature)
        x = self.norm2(some_output)
        return x, queries, top_k_indices

class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len=5000):
        super().__init__()
        position = torch.arange(max_len).unsqueeze(1)
        div_term = torch.exp(torch.arange(0, d_model, 2) * (-math.log(10000.0) / d_model))
        pe = torch.zeros(1, max_len, d_model)
        pe[0, :, 0::2] = torch.sin(position * div_term)
        pe[0, :, 1::2] = torch.cos(position * div_term)
        self.register_buffer('pe', pe)

    def forward(self, x):
        return x + self.pe[:, :x.size(1)]

class SOMETwitter(nn.Module):
    def __init__(self, model_config, some_config):
        super().__init__()
        self.embedding = nn.Embedding(model_config['vocab_size'], model_config['d_model'])
        self.pos_encoder = PositionalEncoding(model_config['d_model'], model_config['seq_len'])
        self.layers = nn.ModuleList([
            SOMETwitterBlock(model_config['d_model'], model_config['num_heads'],
            some_config)

```

```

        for _ in range(model_config['num_layers'])
    ])
self.fc_out = nn.Linear(model_config['d_model'], model_config['vocab_size'])
self.d_model = model_config['d_model']

def forward(self, x, temperature=1.0):
    x = self.embedding(x) * math.sqrt(self.d_model)
    x = self.pos_encoder(x)
    all_queries, all_indices = [], []
    for layer in self.layers:
        x, queries, top_k_indices = layer(x, temperature=temperature)
        all_queries.append(queries)
        all_indices.append(top_k_indices)
    return self.fc_out(x), all_queries, all_indices

@torch.no_grad()
def update_all_keys(self, all_queries, all_indices):
    for i, layer_block in enumerate(self.layers):
        queries = all_queries[i].view(-1, layer_block.some_layer.d_model)
        indices = all_indices[i].view(-1, layer_block.some_layer.top_k)
        layer_block.some_layer.update_keys(queries, indices)

# --- 2. Data Preparation ---

class LanguageModelDataset(Dataset):
    def __init__(self, tokenized_data):
        self.data = tokenized_data
    def __len__(self):
        return len(self.data)
    def __getitem__(self, idx):
        item = self.data[idx]
        inputs = torch.tensor(item['input_ids'])
        targets = inputs.clone()
        targets[:-1] = inputs[1:]
        targets[-1] = -100
        return inputs, targets

def prepare_data(config):
    print("\n--- Part 2: Data Preparation & Configuration ---")
    tokenizer_path = "tinystories-tokenizer-v2.json"
    if not os.path.exists(tokenizer_path):
        print("Training custom tokenizer...")
    dataset = load_dataset("roneneldan/TinyStories", split="train")
    def get_training_corpus():

```

```

        for i in range(0, len(dataset), 1000):
            yield dataset[i : i + 1000]["text"]
    tokenizer_model = Tokenizer(BPE(unk_token="[UNK]"))
    tokenizer_model.pre_tokenizer = Whitespace()
    trainer = BpeTrainer(special_tokens=["[UNK]", "[PAD]", "[EOS]"],
                          vocab_size=config['model']['vocab_size'])
    tokenizer_model.train_from_iterator(get_training_corpus(), trainer=trainer)
    tokenizer_model.save(tokenizer_path)

else:
    print("Tokenizer already exists. Loading from file.")

tokenizer = PreTrainedTokenizerFast(tokenizer_file=tokenizer_path)
tokenizer.add_special_tokens({'pad_token': '[PAD]', 'eos_token': '[EOS]'})
print(f"Custom tokenizer loaded with vocab size: {tokenizer.vocab_size}")

print("\nTokenizing dataset...")
full_dataset = load_dataset("roneneldan/TinyStories", streaming=False)
train_subset = full_dataset['train'].select(range(config['data']['train_subset_size']))
val_subset = full_dataset['validation'].select(range(config['data']['val_subset_size']))

def tokenize_function(examples):
    text_with_eos = [s + tokenizer.eos_token for s in examples["text"]]
    return tokenizer(text_with_eos, truncation=True, padding="max_length",
                    max_length=config['model']['seq_len'], return_tensors="pt")

tokenized_train = train_subset.map(tokenize_function, batched=True,
                                   remove_columns=["text"], num_proc=os.cpu_count())
tokenized_val = val_subset.map(tokenize_function, batched=True,
                               remove_columns=["text"], num_proc=os.cpu_count())

train_dataset = LanguageModelDataset(tokenized_train)
validation_dataset = LanguageModelDataset(tokenized_val)

num_workers = max(2, os.cpu_count() // 2 if os.cpu_count() else 2)
train_loader = DataLoader(train_dataset, batch_size=config['data']['batch_size'],
                         shuffle=True, drop_last=True, num_workers=num_workers,
                         pin_memory=True)
validation_loader = DataLoader(validation_dataset, batch_size=config['data']['batch_size'],
                               drop_last=True, num_workers=num_workers, pin_memory=True)

return train_loader, validation_loader, tokenizer

```

# --- 3. Training & Evaluation Functions ---

```

def calculate_gini(usage_counts):
    counts = usage_counts.cpu().to(torch.float32).numpy()
    if np.sum(counts) == 0: return 0.0
    counts = np.sort(counts)
    n = len(counts)
    index = np.arange(1, n + 1)
    return (np.sum((2 * index - n - 1) * counts)) / (n * np.sum(counts))

def calculate_entropy(usage_counts):
    total_usage = usage_counts.sum()
    if total_usage == 0: return 0.0
    probs = usage_counts / total_usage
    probs = probs[probs > 0]
    return -torch.sum(probs * torch.log2(probs)).item()

def train_epoch(model, dataloader, optimizer, criterion, scheduler, current_temp,
tokenizer_vocab_size):
    model.train()
    total_loss = 0
    scaler = torch.cuda.amp.GradScaler()
    progress_bar = tqdm(dataloader, desc=f"Training (Temp={current_temp:.2f})", leave=False)

    for inputs, targets in progress_bar:
        inputs, targets = inputs.to(device, non_blocking=True), targets.to(device,
non_blocking=True)

        with torch.cuda.amp.autocast():
            # Pass annealing temperature to model
            logits, queries, indices = model(inputs, temperature=current_temp)
            loss = criterion(logits.view(-1, tokenizer_vocab_size), targets.view(-1))

            optimizer.zero_grad(set_to_none=True)
            scaler.scale(loss).backward()
            scaler.unscale_(optimizer)
            torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
            scaler.step(optimizer)
            scaler.update()
            scheduler.step()

        # Update Keys (In-Inference Plasticity)
        model.update_all_keys(queries, indices)

        total_loss += loss.item()
        progress_bar.set_postfix({'loss': f'{loss.item():.4f}', 'lr': f'{scheduler.get_last_lr()[0]:.1e}'})

```

```

    return total_loss / len(dataloader)

def evaluate_epoch(model, dataloader, criterion, tokenizer_vocab_size):
    model.eval()
    total_loss = 0
    progress_bar = tqdm(dataloader, desc="Evaluating", leave=False)

    with torch.no_grad():
        for inputs, targets in progress_bar:
            inputs, targets = inputs.to(device, non_blocking=True), targets.to(device,
non_blocking=True)
            with torch.cuda.amp.autocast():
                # Sharpen during eval (low temp)
                logits, _, _ = model(inputs, temperature=0.5)
                loss = criterion(logits.view(-1, tokenizer_vocab_size), targets.view(-1))
                total_loss += loss.item()
            progress_bar.set_postfix({'loss': f'{loss.item():.4f}'})

    return total_loss / len(dataloader)

def plot_losses(train_losses, val_losses, config):
    epochs = len(train_losses)
    plt.figure(figsize=(10, 6))
    plt.plot(range(1, epochs + 1), train_losses, 'b-o', label='Training Loss')
    plt.plot(range(1, epochs + 1), val_losses, 'r-o', label='Validation Loss')
    title = f"SoME v4.0 Run: {config['run_name']}"
    plt.title(title)
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)
    plt.xticks(range(1, epochs + 1))
    filename = f"loss_curve_{config['run_name']}.png"
    plt.savefig(filename)
    plt.show()

```

# --- 4. Main Execution Function ---

```

def main(config):
    print(f"\n--- Starting Experiment: {config['run_name']} ---")

    # 1. Data
    train_loader, val_loader, tokenizer = prepare_data(config)

```

```

# 2. Model Initialization
print("\n--- Part 3: Model Definition ---")
model = SOMETransformer(config['model'], config['some_layer']).to(device)

if hasattr(torch, 'compile'):
    print("\nCompiling the model for faster training...")
    model = torch.compile(model)

# 3. Training Setup
print("\n--- Part 4: Training, Evaluation, and Metrics ---")
# Only optimize parameters that require grad (Router/Keys, NOT Experts)
optimizer = torch.optim.AdamW([p for p in model.parameters() if p.requires_grad],
                             lr=config['training']['learning_rate'],
                             betas=(0.9, 0.95), weight_decay=0.1)

criterion = nn.CrossEntropyLoss(ignore_index=-100)
total_steps = len(train_loader) * config['training']['num_epochs']
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=total_steps)

total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"\nTotal parameters: {total_params/1e6:.2f}M")
print(f"Trainable parameters: {trainable_params/1e6:.2f}M ({100 * trainable_params / total_params:.2f}%)")
print(f"Total training steps: {total_steps}")

# 4. Training Loop with Annealing
train_losses, val_losses = [], []
best_val_loss = float('inf')
model_save_path = f"best_model_{config['run_name']}.pth"

# Temperature Schedule: High Entropy (2.0) -> Low Entropy (Configured)
start_temp = 2.0
end_temp = config['training']['training_temp']

for epoch in range(config['training']['num_epochs']):
    print(f"\n--- Epoch {epoch+1}/{config['training']['num_epochs']} ---")

    # Calculate current temperature
    progress = epoch / config['training']['num_epochs']
    current_temp = start_temp - (start_temp - end_temp) * progress
    print(f"Current Router Temperature: {current_temp:.4f}")

```

```

train_loss = train_epoch(model, train_loader, optimizer, criterion, scheduler,
                        current_temp, tokenizer.vocab_size)

val_loss = evaluate_epoch(model, val_loader, criterion, tokenizer.vocab_size)
perplexity = math.exp(val_loss)

train_losses.append(train_loss)
val_losses.append(val_loss)

# Metrics Inspection (Middle Layer)
model_to_inspect = model._orig_mod if hasattr(model, '_orig_mod') else model
mid_layer = model_to_inspect.layers[config['model']['num_layers'] // 2].some_layer
usage_counts = mid_layer.usage_count
gini_coeff = calculate_gini(usage_counts)
entropy_val = calculate_entropy(usage_counts)

# Check Phoenix Activity
dead_experts = (mid_layer.key_store.norm(dim=-1) <
mid_layer.spawn_threshold).sum().item()

print(f"Epoch {epoch+1}: Train Loss = {train_loss:.4f}, Val Loss = {val_loss:.4f}, Val Ppl = {perplexity:.2f}")
print(f" Middle Layer Metrics: Gini = {gini_coeff:.3f}, Entropy = {entropy_val:.3f}, Dead Experts Pending Respawn: {dead_experts}")

if val_loss < best_val_loss:
    best_val_loss = val_loss
    torch.save(model_to_inspect.state_dict(), model_save_path)
    print(f" Model saved as {model_save_path}")

# 5. Finalization
print("\n--- Training Complete for {config['run_name']} ---")
plot_losses(train_losses, val_losses, config)

```

# Cell 3: Experiment Configuration & Execution

```

# --- Define the configuration for the experiment ---
config = {
    "run_name": "v4_C2_WidthFix_Phoenix",

    "data": {
        "train_subset_size": 10000,
        "val_subset_size": 1000,
        "batch_size": 32, # Adjusted for C2 scale

```

```
},  
  
"model": {  
    "vocab_size": 8192,  
    # C2 Dimensions: Width 768 caused collapse in v3  
    "d_model": 768,  
    "num_heads": 12,  
    "num_layers": 8,  
    "seq_len": 768,  
},  
  
"some_layer": {  
    "num_experts": 128,  
    "d_ffn": 1536,  
    "top_k": 4,  
    "init_method": "sparse",  
  
    # v4 Heuristics  
    "alpha": 0.015,  
    "beta": 0.001,  
    "delta": 0.001,  
    "respawn_threshold": 0.1, # <--- NEW: Expert death threshold  
  
    "theta_percentile": 0.05,  
    "warmup_steps": 400,  
    "ema_decay": 0.995,  
  
    "ablation_flags": {  
        "use_alpha": True,  
        "use_beta": True,  
        "use_delta": True  
    }  
},  
  
"training": {  
    "num_epochs": 2,  
    "learning_rate": 6e-4,  
    "training_temp": 1.0, # Annealing will go 2.0 -> 1.0  
}  
}  
  
# --- Run the experiment ---  
main(config)
```

---

```

# Cell 4: Multi-Layer SoME Analysis & Diagnostics Dashboard v1.3
# RUN THIS CELL AFTER A TRAINING RUN

print("\n--- Part 1: Dashboard Setup ---")
import umap
import seaborn as sns

# Ensure the latest configuration is loaded
model_path_to_load = f"best_model_{config['run_name']}.pth"
tokenizer_path = "tinystories-tokenizer-v2.json"

if os.path.exists(model_path_to_load) and os.path.exists(tokenizer_path):
    print(f"Loading best model from: {model_path_to_load}")
    tokenizer = PreTrainedTokenizerFast(tokenizer_file=tokenizer_path)
    tokenizer.add_special_tokens({'pad_token': '[PAD]', 'eos_token': '[EOS]'})

analysis_model = SOMETransformer(config['model'], config['some_layer']).to(device)
analysis_model.load_state_dict(torch.load(model_path_to_load))
analysis_model.eval()

# --- Aggregate Utilization Analysis ---
middle_layer_idx = config['model']['num_layers'] // 2
middle_layer = analysis_model.layers[middle_layer_idx].some_layer
usage_counts = middle_layer.usage_count.cpu()

plt.figure(figsize=(12, 6))
plt.hist(usage_counts.numpy(), bins=50, color='skyblue', edgecolor='black')
plt.yscale('log')
plt.title(f"Expert Usage Counts ({config['run_name']} - Layer {middle_layer_idx})")
plt.show()

# --- Key Store UMAP ---
print("Running UMAP projection...")
key_store_data = middle_layer.key_store.cpu().numpy()
reducer = umap.UMAP(n_neighbors=15, min_dist=0.1, n_components=2, metric='cosine')
embedding = reducer.fit_transform(key_store_data)

plt.figure(figsize=(12, 10))
plt.scatter(embedding[:, 0], embedding[:, 1], c=usage_counts, cmap='viridis', s=20, alpha=0.7)
plt.colorbar(label='Usage Count')
plt.title(f"UMAP of Expert Keys (Color=Usage)")
plt.show()
else:
    print("Model file not found. Run training first.")

```

---